StratusLab

Enhancing Grid Infrastructures with Virtualization and Cloud Technologies

First Year Software Integration Report

Deliverable D4.3 (V1.0)

15 June 2011

Abstract

This document reports on the integration activities performed during the first year of StratusLab, focusing on lessons learned and identifying the areas of improvements, not only for the first release of the StratusLab software, but also on the process used during the development, integration and test, such that the second year of the project builds on these lessons and improves its performance.

The information contained in this document represents the views of the copyright holders as of the date such views are published.

Contributors

Name	Partner	Sections
Marc-Elian Bégin	SixSq	All
Louise Merifield	SixSq	All

Document History

Version	Date	Comment
0.1	2 June 2011	Initial version for comment.
0.2	12 June 2011	Corrections from Louise Merifield.
0.3	14 June 2011	Added execution summary.
1.0	15 June 2011	Final version.

Contents

List of Figures

List of Tables

1 Executive Summary

StratusLab is a two year project, with a challenging programme of work. The project exists in the context of cloud and grid computing, a fast and very dynamic field in distributed software engineering. In order to manage this challenging situation, the project decided to adopt an agile software development process, called Scrum, alongside a set of engineering practices, some taken from eXtreme Programming (XP).

This report describes the strategy and techniques used during the first year of the project to deliver this architecture, over a set of incremental deliveries, building on the strength of the previous, generating feedback from users, and feeding this back into the prioritisation of the functionality released with every iteration.

While the Agile eco-system is rich and diverse (e.g. Extreme Programming - XP, Scrum, Lean Development, Kanban, Core Principals), StratusLab chose to focus its organisational structure on Scrum. Scrum has a clear model and is prescriptive of the dynamics between the customer and the provider, the emergence of functional specification (i.e. Product Backlog) and the rhythm of development (i.e. sprints). While other methods would also apply, such as XP (as well as derivatives like Test-driven development - TDD), we focused our effort on the Scrum part of the development process.

The Scrum process described early in this document focuses on the management of functionality and prioritization of features for each StratusLab release. This report then changes its focus towards the engineering practices used in the integration and test activities required to deliver high-quality StratusLab software.

To deliver a quality solution the StratusLab project employs a set of practices of continuous automated integration and deployment in the course of the development process. The set of practices (properly pipelined, automated and frequently executed) allows for rapid feedback on the quality of the software stack at all levels as well as its delivery models.

The project uses Git (a distributed revision control system) as its code repository. The build procedures themselves in StratusLab are implemented using Maven2. This tool is the de-facto standard in the Java world, but also provides support for other languages. Maven2 provides a rich set of plugins, capable of performing complex tasks.

In order to guarantee that everything is working as we move forward with new features, improvements and fixes, the project uses a continuous integration (CI)

infrastructure, composed of a main server and a number of dedicated 'slave' machines.

Having reviewed and explained the methodologies and processes used to integrate and test StratusLab, the last part of this report brings together important lessons learned. These include a wide range of lessons:

- Further improvements to the continuous integration system

- Documentation updates for every development and integration task

- Further automation of machine re-imaging to ensure clean installations using Quattor

- Push patches back to external code providers faster to reduce flakiness in our build system

- Continue working on HEAD/master branch whenever possible

- Continue to commit on a regular basis

- Continue to encourage a stop-the-line culture regarding Hudson jobs to avoid committing code on a broken system

- Improve and increase variety of modern operating systems support for StratusLab user tools and cloud system.

- Continue to support manual and automated installation capabilities

The first year of the project was very productive with several releases and numerous sprints delivering significant functionality towards the project goal of delivering a fully functional high quality cloud distribution. The build and test activities summarised in this report were also a great source of lessons learned, which we are convinced will help the project improve even further in its second year.

2 Introduction

StratusLab is a two year project with a challenging programme of work. The project exists in the context of cloud and grid computing, a fast and very dynamic field in distributed software engineering. In order to manage this challenging situation, the project decided to adopt an agile software development process, called Scrum, alongside a set of engineering practices, some taken from eXtreme Programming (XP).

Figure 2.1 shows the high-level architectural vision of StratusLab v1.0. This report describes the strategy and techniques used during the first year of the project to deliver this architecture, over a set of incremental deliveries, building on the strength of previous experience, generating feedback from users, and feeding this back into the prioritisation of the functionality released with every iteration.

As you can see in this diagram, StratusLab is composed of a number of components and services, each integrated and tested individually and as a whole, using an automated procedure and on a regular basis (i.e. generally several times per day). This report describes how the development, integration and test activities are orchestrated and synchronised to produce the StratusLab distribution.
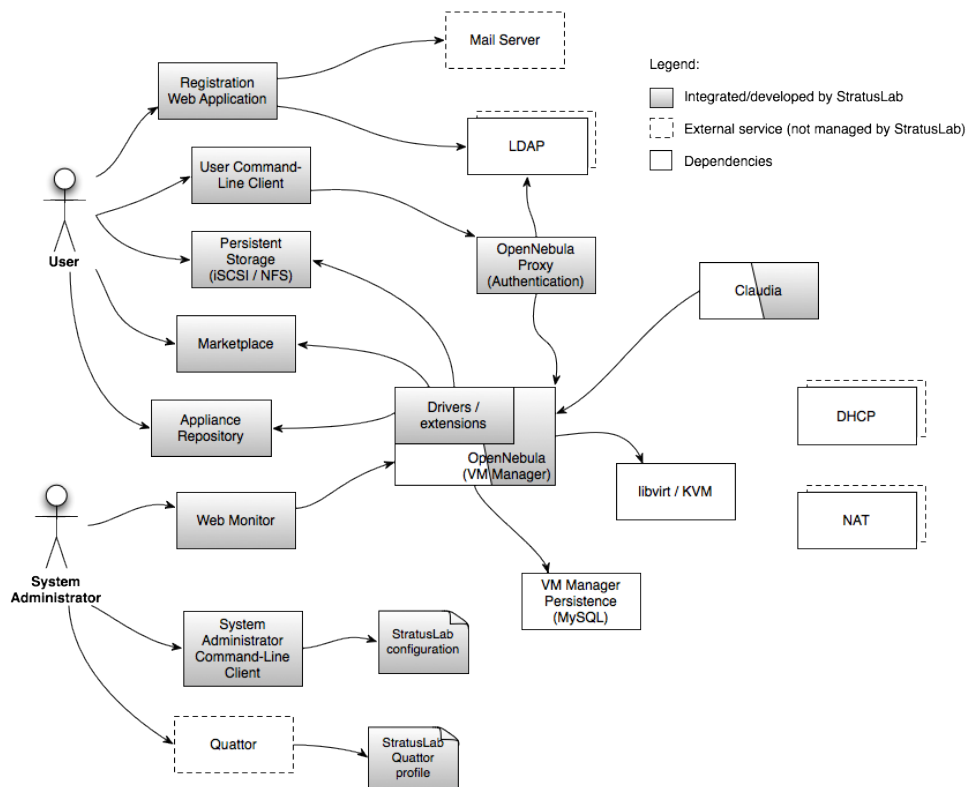
Figure 2.1: *StratusLab v1.0 Architectural Vision*

3 Agile Process: Scrum

3.1 Agile Overview

While the Agile eco-system is rich and diverse (e.g. Extreme Programming - XP, Scrum, Lean Development, Kanban, Core Principals), StratusLab chose to focus on organisational structure on Scrum. Scrum has a clear model and is prescriptive of the dynamics between the customer and the provider, the emergence of functional specification (i.e. Product Backlog) and the rhythm of development (i.e. sprints). While other methods would also apply, such as XP (as well as derivatives like Test-driven development - TDD), we focused our efforts on the Scrum part of the development process.

The main reasons for choosing Scrum were:

- Client centric approach encouraging convergence towards fulfilling the real needs

- Improved project visibility

- Promotes trust between all stakeholders

- Controlled incorporation of changes

- Continuous inspection and improvement of the process

- Higher quality software

- Simpler solutions, easier to maintain and evolve

- More predictable delivery of functionality

While a complete description of Agile and Scrum is outside the scope of this report, the following explains the main artefacts and processes involved in Stratus-Lab's implementation of Scrum. In Scrum, each iteration (called a sprint) starts with a planning meeting. During this meeting, high-priority requirements are reviewed, analysed, decomposed into tasks and selected for the next sprint. The duration of each sprint varies normally between one and four weeks, with a trend towards shorter sprints. StratusLab decided to standardise at around three weeks, to take into account the integration overhead and the distributed nature of our teams.

Figure 3.1: *Scrum Overview*

Each sprint ends with a sprint review, composed of two events: a retrospective and a demo. The objective of the retrospective is to review the past sprints' performance, with a focus on emulating what was particularly effective, while eliminating impediments. For the demo, the team assembles and deploys the software and presents the implementation of the requirements selected for that sprint. Different stakeholders can be invited to this event, if necessary, including for example endusers. This is an important generator of feedback and insights, which can be fed into the requirements, and provide new data for prioritisation.

Another important Scrum event (but also required by most agile methods) is the daily meeting, daily stand-up or daily Scrum. The sole purpose of this meeting, which should not exceed 15 minutes, is to foster continuous and fluid communication between all teams. Stand-up meetings focus on reporting what has been accomplished since the last meeting, what is planned until the next meeting and identification of any impediments. Further, to alleviate the fact that for StratusLab, stand-up meetings cannot be performed face-to-face, a new item to the standard topics is a mention of the state of the Hudson continuous integration server, such that any failures are highlighted to all and corrective actions are immediately scheduled. StratusLab stand-up meetings take place Monday to Friday at 10:30 sharp, Paris time.

Scrum also defines three key roles:

Product Owner The person responsible for maintaining the Product Backlog by representing the interests of the stakeholders.

Scrum Master The person responsible for the Scrum process, making sure it is used correctly and maximizing its benefits.

The Team A cross-functional group of people responsible for managing itself to develop the product.

These three roles are mapped differently to each project structure, depending on the project, its stakeholders, contractual setup, distribution of skills and locations across participants, etc. In StratusLab, the role of Product Owner is fulfilled by a group composed of all work-package leaders and senior technical members of the project. Charles Loomis, the project director, arbitrates and help focus the group regarding prioritisation. The role of Scrum Master is generally held by Marc-Elian Bégin, WP4 leader, but is regularly filled by other members when required.

Each sprint is concluded with a demo. This is the opportunity to show the project members, via a alive demo, the completion of each user story, improvements and bug fixes. Over the course of the project, the definition attached to completing an item of work (definition of 'done') evolved and is expected to continue evolving. This natural evolution also follows the maturity of our build and test infrastructure, discussed in more detail in Chapter 4.5. This improvement to the doneness of our tasks is important the reduce the risk that problems are discovered during deployment by WP5, or worse that we break a feature already delivered in a previous version (also called regression).

An additional reason for choosing relatively short three-week sprints is that, assuming we can release often enough, bug fixes can be released with normal releases, as opposed to requiring special bug fixing releases. While this policy might have to be revisited if critical problems are detected to which an urgent fix is required, it offers a much simpler rollout model.

3.2 StratusLab First Year Scrum Measurements

The main artefact driving the Scrum process is the 'Product Backlog'. For StratusLab, the Product Backlog is captured using JIRA (and the GreenHopper plugin).

While the planning meetings, demos and daily stand-up events, described in the previous section, provide the 'tactical' process for efficient project execution, we also need a longer term and strategically focused think tank. This important function is fulfilled by the Technical and Scientific Coordination Group (TSCG), chaired by Ruben Montero from UCM, to which all work-package leaders and senior project members contribute. From the set of priorities defined by the TSCG, the product backlog is maintained (i.e. items added and removed), including clear priorities.

Figure 3.2 illustrates the virtuous cycle feeding the StratusLab agile process. The diagram clearly shows the important role and actions each activity has and performs. This feeds into the feature prioritisation performed by the TSCG, as well as the architectural vision required to provide a solution. Using our agile approach we can then more easily follow and adapt to deliver these features.

As mentioned above, the StratusLab Product Backlog is captured with JIRA using three types of tickets:

User Story A feature requiring several tasks to complete, possibly involving several partners, but sized such that it can be implemented within a single sprint.
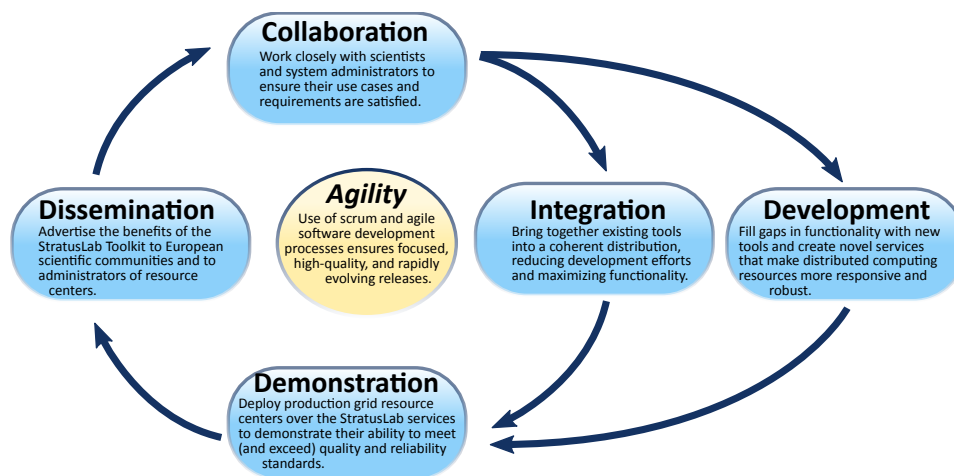
Figure 3.2: *Work-Package Interactions*

Table 3.1: *First Year StratusLab Releases*

Release Number	Date	Sprint
v0.1	9 November 2010	Sprint 5
v0.2	23 December 2010	Sprint 7
v0.3	15 March 2011	Sprint 10

Bug A bug, requiring a simple fix, normally affecting a single component

Improvement A small upgrade in current functionality, not requiring a new User Story

This level of granularity allows us to manage all items during sprint execution.

During the first year of the project, we have completed 14 sprints (including sprint 0). Figure 3.3 gives an overview of the number of items completed for each sprint. We can also correlate important project events on this picture. For example, the dip in sprints 7 and 8 correspond to the winter and new year break, where several team members were on holiday, therefore reducing the task force available to complete items. The surge in bugs in sprint 3 corresponds to our first release (v0.1) where certification in view of production deployment identified issues. From sprint 10, the number of 'improvements' started to grow, corresponding to feedback from usage, expressed as small improvement requests, as opposed to new functionality.

During the first year of the project, three releases (see details in Table 3.1 were performed, with two more in June 2011 (i.e. v0.4 and v1.0).

Each release was tested and deployed on dedicated test machines, prior to being released to WP5.

Figure 3.3: *Sprints Overview*

4 Engineering Practices

The Scrum process described in Chapter 3 focuses on the management of functionality and prioritization of features for each StratusLab release. This chapter focuses on the engineering practices used in the integration and test activities required to deliver high-quality StratusLab software.

4.1 Continuous Integration and Deployment

To deliver a quality solution the StratusLab project employs a set of practices of continuous automated integration and deployment in the course of the development process. The set of practices (properly pipelined, automated and frequently executed) allows for rapid feedback on the quality of the software stack on all of its levels as well as its delivery models.

To facilitate continuous integration effort the StratusLab project uses Hudson.

Build Automation Build automation is handled using Apache Maven. All separate software components of the StratusLab project (identified as separate projects in the git repository) are Maven projects. This allows for a common build interface and specification of concise instructions for the components testing, building and upload to the project's distribution repositories.

Building all commits to baseline To reduce the number of conflicting code changes, regular commits (many times a day by a single developer) to the code baseline are encouraged and being practiced. The commits are immediately unit-tested, built, deployed and functionally tested to reduce the window between commit and the feature/fix actual functional usage/testing. The latter facilitates an immediate awareness by the developers of any possible failures caused by the code changes. In case of failures, responsible people are notified immediately by email.

Deployment, Integration and System Testing A special infrastructure comprising physical nodes as well as virtual machines was deployed by the project to resemble as much as possible a scaled production environment. The infrastructure is used to exercise per-component feature testing and what is most important run integration as well as system tests by simulation of user-like system usage on dynamically deployed latest snapshot versions of the software stack.

Release Automation To ease the software stack release procedure special Maven build targets were defined for each component and per-component release jobs were created in Hudson. Also, a common catch-all release job per supported Linux distribution triggering the per-component release jobs was created. The latter allows the release of the StratusLab software stack with virtually a single click.

4.2 Configuration Management and Version Control

The project uses Git (a distributed revision control system) as its code repository. The project's Git repository[1] contains the following logically identified sub-projects each often composed of more than one sub-module, from which the build procedure produces a number of software packages.

stratuslab-authn Authentication proxy for OpenNebula and Claudia, providing support for local or LDAP managed username/password, as well as X509 and grid certificates

stratuslab-benchmarks A set of standard benchmarks for the cloud installation

stratuslab-claudia Claudia system

stratuslab-client End-user client for remote creation and management of virtual machines, system administrator installation and configuration tools, web-monitor for simple monitoring of physical nodes and virtual machines.

stratuslab-image-recipes Image recipes for creating base virtual images from standard operating systems

stratuslab-marketplace Marketplace providing a service for query and registration of virtual appliances metadata (both machine and disk images), with cryptography support of authors and endorsers of the images.

stratuslab-one StratusLab custom OpenNebula build, integrating the latest upgrades, patches and customisation for StratusLab

stratuslab-quattor Set of Quattor templates for automated installation and configuration of StratusLab

stratuslab-registration A registration web application providing confirmation workflow for users of a specific StratusLab installation.

stratuslab-storage A persistent storage web application for the creation and management of persistent disk for virtual image instances.

project-documents All project documents, including deliverables and dissemination material

---

[1] Based at https://code.stratuslab.eu/git/

Using a common Git store, all code contributors can share code. Like all good version control systems, Git provides a merge feature, ensuring that conflicting commits can be resolved without loss of information. All team members are encouraged to commit (and push) their code often.

An important component in StratusLab is OpenNebula. In order to better support StratusLab, the OpenNebula team at UCM developed new features and provided bug fixes requested and reported by StratusLab integration and operation activities. However, OpenNebula has a different and much slower release cycle, incompatible with StratusLab's aggressive release strategy. In order to resolve this issue the UCM team created a clone Git repository (OpenNebula also uses Git for managing its code) dedicated to StrartusLab. This allows UCM to rollout directly to this repository code so that StratusLab can release, independently from OpenNebula releases. In order to avoid the risk of divergence between the master OpenNebula code base and StratusLab's custom version, UCM regularly merges back to the master code base.

The same applies for patches. StratusLab team members regularly modify OpenNebula extensions (e.g. to fix bugs, extend functionality). This is done by defining patch files that are applied by the Maven build procedure (see next section for detail). These patches have the advantage of providing immediate fixes, but over time can become expensive to maintain as the code they patch changes over time. To mitigate this situation, OpenNebula regularly integrates these patches into the mainstream code base, so that these patches can be removed.

4.3 Building Software

The build procedures themselves in StratusLab are implemented using Maven2. This tool is the de-facto standard in the Java world, but also provides support for other languages. Maven2 provides a rich set of plugins, capable of performing complex tasks. Here are the operations that we execute using Maven2 when building StratusLab packages:

- extract and install dependencies

- build binaries

- execute unit tests

- package

- install

- configure

- deploy to remote server

- release
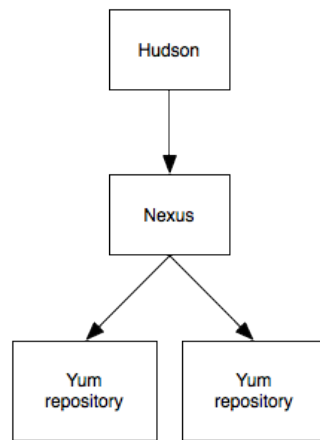
Figure 4.1: *Packaging Publishing*

While the details of every step listed above goes beyond the scope of the current report, here are a few important aspects worth mentioning regarding these steps. The packages are currently only created in the RPM format. While we have partial support for Debian packages (compatible with Ubuntu), we are not currently able to release a full distribution in .deb package format.

Having said that, we also release the end-user command-line client in tarball format, which is tested on CentOS, Fedora 14, Ubuntu and Windows.

Maven2 generates the packages and deploys them to a Maven2 server-side component called Nexus. The Nexus server provides a convenient way to provide dependencies. This is important so that developers can very quickly create a development environment using Maven without having to install software from different sources.

Maven is also used to generate binary packages (e.g. RPM) which are published and uploaded by the deploy target to Nexus. From there, RPMs are then registered to package repositories, such as YUM for RPM. This process is illustrated in Figure 4.1. The packages generated by Maven are also configured to include any runtime dependencies, such that the packaging system can pull all required dependencies during installation.

Further, the same process is used to handle snapshot and release packages. This means that we can have a clear separation between the packages that are produced with every build (i.e. snapshots), and the packages that are released and therefore visible by our users.

For release packaging, the Maven2 release plugin is used, which automatically adjusts the version number of every package and incrementing the number, with Git commits at every stage, such that the process is reversible and fully traceable.

4.4 Hudson

In order to guarantee that everything is working as we move forward with new features, improvements and fixes, the project uses a continuous integration (CI) infrastructure, composed of a main server and a number of dedicated 'slave' machines.

The project chose Hudson as the CI server[2]. Hudson is open source and the de-facto standard CI solution at the moment. Hudson executes a number of 'jobs', from triggers ranging from code commits in Git, on a schedule, or from the successful completion of a parent job. The jobs created and pipelined comprise code checkout, unit-testing, build, deployment, functional multi-service testing as well as release of all the StratusLab services and components on the supported platforms.

At the time of writing there are 61 registered jobs with 43 of them active. Currently, the majority of the disabled jobs are the ones that were created for CentOS 5.5 Linux distribution, support for which was dropped by the project in a favor of Fedora 14. We have decided to keep the old CentOS jobs to ease the future support of CentOS 6 by simple re-profiling of them, when CentOS 6 is released.

For example, every day, we test a complete installation of the latest snapshot version of StratusLab. At 2:00 in the morning (Paris time), two machines are re-imaged with a pristine operating system installation. This re-imaging process is managed by Quattor at LAL. At 3:00, a Hudson scheduled job triggers and executes the job workflow illustrated in Figure 4.2. After the re-image, the machine is upgraded such that it contains the right dependencies, as specified in the online documentation. Then the StratusLab front-end (including OpenNebula) is installed and configured. From that point, several jobs are executed, testing features and/or adding supplementary software before testing it. For this process to work, two Quattor controlled machines are used, one for the front-end, the other for the node.

Another important workflow is the building of each package. This process is illustrated in Figure 4.3, in this case for releasing the packages. From a single Hudson job trigger, we are now able to preform a release of the complete StratusLab distribution. Each package can also be released individually, by directly triggering the corresponding job.

We currently have a significant Hudson infrastructure, comprising of six physical servers and two virtual machines. We try to use as many virtual machines as possible, instead of physical servers. However, since several of the Hudson jobs use and test virtualisation, we require a physical server, since for example it is not possible to run a virtual machine inside another virtual machine (unless para-virtualisation is used inside full-virtualisation, which adds its own complexity and biases).

We did not construct our build and test infrastructure overnight. On the contrary, we have built it over successive sprints, growing it as needs were identified,

---

[2]Deployed at http://hudson.stratuslab.eu:8080/

such that we could release software very early after project kickoff, carefully trading off effort invested in the build and test infrastructure and tools from integration and development tasks.

The Hudson server will be polling the Git repository regularly, checking for recent commits. Then, the build job will be triggered, generating snapshot packages, following by triggering the job labelled 'cloud Install NFS' in Figure 4.2. This means that developers receive feedback a few minutes after each commit, if their change has introduced a regression. This near immediate feedback gives the developer an opportunity to fix the problem while the problem domain is fresh in his/her mind, as opposed to being told that a problem occurred following a change committed days or even weeks before. Finally, assuming that all tests were green (passed) before the commit, the sources of errors are much narrower, compared with integration activities trying to integrate several changes at once.

4.5 Installation Strategies

StratusLab support two installation strategies:

Manual Installation This method supports Fedora14 and similar systems. It consists of a set of commands that perform the necessary installation steps based on information in a single configuration file.

Quattor Installation This method supports an automated installation of Fedora14 machines, which can scale to very large deployments. This method is supported by Quattor, for which StratusLab maintains specific profiles available in the standard StratusLab distribution.

Both methods are used in our automated build and test procedures, such that they are both tested regularly in order to identify any discrepancies that could occur between the two systems. This ensures that whichever the method chosen by system administrator, the resulting StratusLab deployment will be nearly identical.
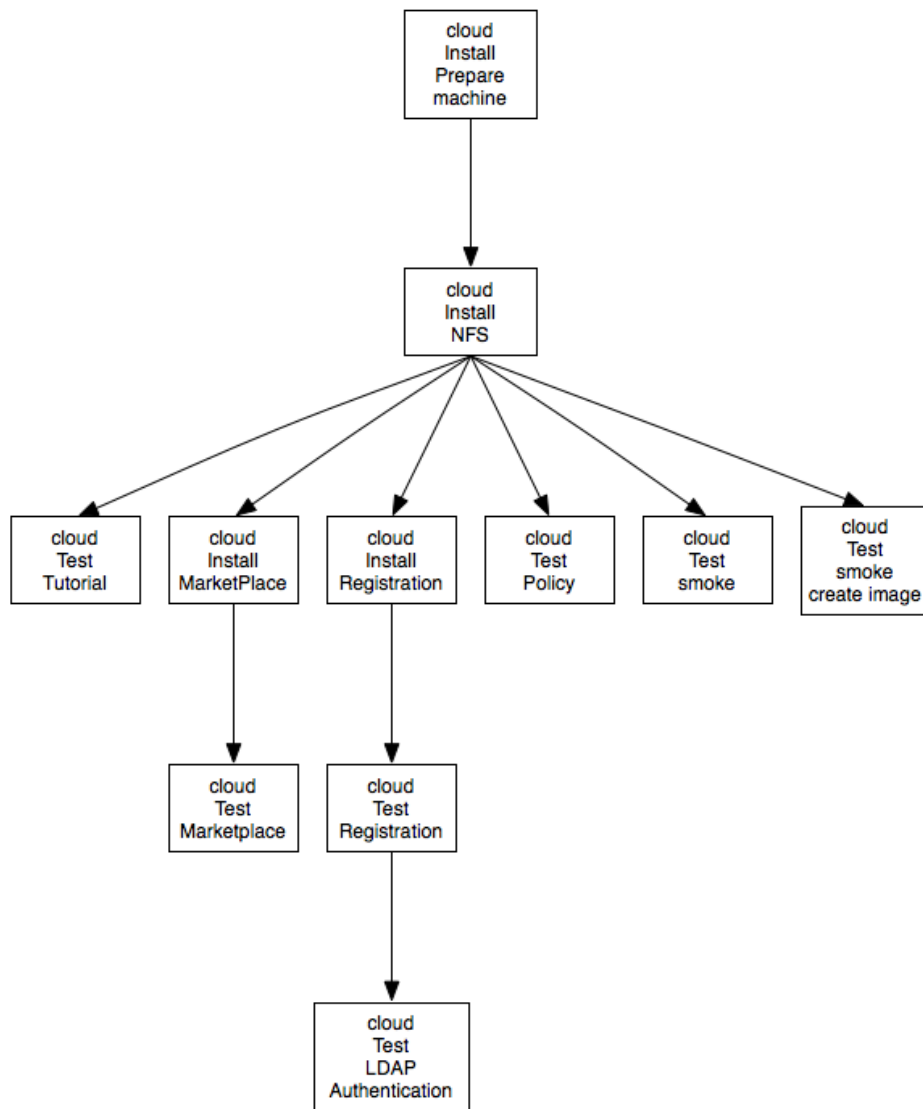
Figure 4.2: *Hudson Jobs for Daily Smoke Tests on Full Installation*

Figure 4.3: *Hudson Release Jobs*

5 Lessons Learned

Having reviewed the agile process the project uses to manage requirements and development, followed by the engineering practices used to develop, integrate and test the StratusLab software, we will now summarise lessons learned, with which we hope to improve during the second year of the project.

1. Earlier this year we lost the main Hudson server, due to a hardware failure. We took advantage of this event to reshape the Hudson server and redefine the job topology, making it more flexible and efficient. In the future however, not only the server should be backed up (as it is now), we should also put the Hudson configuration files (e.g. main configuration and job definitions) under Git for version control. This would allow us to roll back to previous configuration more easily if and when required.

2. Despite our best efforts, each production system upgrade was not completely straightforward, with longer than expected down time. We believe that releasing more often, closer to sprint cycles, would improve our ability to better upgrade, with minimal disruption for users, while providing new features more often. As more sites deploy StratusLab however, we have to be careful not to alienate system administrators with too frequent releases. Further, we need to ensure that upgrade paths are tested and solid, such that no loss of data occurs when upgrading to a new release. Finally, we could consider releasing different components are different rates, for example releasing end-user client more often to provide them with new features more regularly.

3. Another aspect with which we must make a conscious effort is to ensure that documentation is updated as we go along and that this update is added to the definition of done when completing any relevant task. This is important for all stakeholders of the project, such that documentation is available to support integration and test activities, but also for (pre-)production deployment and user evaluation.

4. The introduction of Quattor controlled re-imaging of test machines had a significant positive impact on stabilising and improving the code quality, test coverage and reliability of our installation and configuration tools. As we move forward and support more operating systems, it will be critical that

Quattor controlled machines are provided for all of these, such that we can extend our current testing strategy to these.

5. Our quick patching mechanism works, but must be integrated on a regular basis by the developers of the patched system to avoid escalation in patch maintenance effort. This is now working very well for OpenNebula, and we need to continue ensuring that patches found in the project are considered for integration directly in the original code base, and removed as patches.

6. Another good decision was the creation of a dedicated OpenNebula Git repository for StratusLab, with active merging back to the main OpenNebula release branch. This provides the right balance between being able to provide StratusLab with custom OpenNebula features early, without having to wait for an official OpenNebula release, while being able to benefit from advancements from OpenNebula mainstream development.

7. Our policy of encouraging development on the master branch (or HEAD in other version control system) is paying off. The classic alternative is for each development to take place in separate branches and merging the branches when ready for integration or certification. This can cause headaches and conflicts which can be a significant source of risk and hidden delays. Since our sprints are relatively short and user stories specifically designed to be short and focused, it means that we can afford to develop directly on the master branch, where each commit is automatically built and tested by our Hudson system, including function, system and end-to-end tests.

8. A cousin to this technique of working on the master branch is to commit (and push) often. This requires discipline, where simple and small steps are made when developing, yet this is largely done by all teams, with few exceptions. This technique also has the advantage of being able to back track quickly when jobs turn red after committing faulty code and makes troubleshooting simpler.

9. As more features were integrated and the number of services and components grew in StratusLab, so did the number of jobs in Hudson, as well as the number of resources required to run these. During the third quarter, several jobs remained red (broken) for several days, or were not executed for several days, while team members were busy working on their sprint tasks. This meant that the work required to bring all jobs back to green was long, difficult and often boring. To address this issue, during quarter four, we decided to introduce a Lean technique called 'stop the line'. This meant that we agreed that as soon as a job failed, the person or team responsible for this job would stop their current activity and fix the problem. Within a single sprint, the Hudson server became a lot more stable, a red is now the exception instead of the rule, which was becoming too often the case. We need to ensure that this reflex is maintained during the second year.

10. During sprint 12, we decided to switch our baseline operating system from CentOS 5.5 to Fedora 14. This was a big move, requiring changes to both software and infrastructure. Our ability to deliver this transition within a single sprint is remarkable and probably shows good command of our code base and its dependencies on runtime environment and operating systems. This should make supporting more operating systems relatively straightforward, assuming that our build and test tooling can follow.

11. Continue to support manual and automated installation capabilities is an important asset of the project, giving choices to system administrators willing to deploy StratusLab on their resources. The effort put in to synchronise the behaviour of the two system is also paying-off such that both systems use the same assumptions resulting in almost identical setups.

6 Conclusion

The first year of the project was very productive with several releases and numerous sprints delivering significant functionality towards the project goal of delivering a fully functional high quality cloud distribution. The build and test activities summarised in this report are also a great source of lessons learned, lessons we are convinced will help the project improve further in its second year.