



Enhancing Grid Infrastructures with
Virtualization and Cloud Technologies

Operation of Site Running StratusLab toolkit v1.0

Milestone MS11 (V1.0)
4 April 2011

Abstract

This document reports on one of the significant project milestones concerning the operation of a cloud site running StratusLab's cloud distribution (or toolkit). It is currently running v0.3 of the StratusLab distribution—a beta release of the upcoming 1.0 production release. In particular it presents the WP5 reference cloud service, the evolution of the service up to Month 10 of the project, and the way this service has been exploited to date. A major milestone regarding the exploitation has been the deployment of the first virtualized production grid site running on the project's reference cloud service.



StratusLab is co-funded by the
European Community's Seventh
Framework Programme (Capacities)
Grant Agreement INFSO-RI-261552.



The information contained in this document represents the views of the copyright holders as of the date such views are published.

THE INFORMATION CONTAINED IN THIS DOCUMENT IS PROVIDED BY THE COPYRIGHT HOLDERS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE MEMBERS OF THE STRATUSLAB COLLABORATION, INCLUDING THE COPYRIGHT HOLDERS, OR THE EUROPEAN COMMISSION BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THE INFORMATION CONTAINED IN THIS DOCUMENT, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

Copyright © 2011, Members of the StratusLab collaboration: Centre National de la Recherche Scientifique, Universidad Complutense de Madrid, Greek Research and Technology Network S.A., SixSq Sàrl, Telefónica Investigación y Desarrollo SA, and The Provost Fellows and Scholars of the College of the Holy and Undivided Trinity of Queen Elizabeth Near Dublin.

This work is licensed under a Creative Commons Attribution 3.0 Unported License
<http://creativecommons.org/licenses/by/3.0/>



Contributors

Name	Partner	Sections
Evangelos Floros	GRNET	All
Charles Loomis	LAL/CNRS	Formatting and proofing
Christophe Blanchet	IBCP/CNRS	Bioinformatics Application
David O'Callaghan	TCD	Internal review

Document History

Version	Date	Comment
0.1	21 Mar. 2011	First draft for internal review.
0.2	23 Mar. 2011	Release candidate version
1.0	04 Apr. 2011	Final version

Contents

List of Figures	5
List of Tables	6
1 Introduction	7
2 Reference Cloud Service	8
2.1 Service Deployment Evolution	8
2.2 Physical Configuration.	8
2.3 Service Provisioning.	9
2.4 Site Monitoring	10
2.5 Sample Applications.	10
3 Hosting of Grid Services	12
References	13

List of Figures

2.1	StratusLab's physical infrastructure in GRNET's datacenter. . . .	9
2.2	Web monitoring of cloud site (hosting nodes)	10

List of Tables

2.1	Summary of the physical resources of the cloud site.	8
-----	--	---

1 Introduction

The StratusLab project is working on bridging grid and cloud computing technologies. This entails both the provision of grid services on top of IaaS (Infrastructure as a Service) cloud infrastructures as well as the expansion of the cloud access paradigms to incorporate grid computing attributes like collaboration and resource sharing. The main product of this effort is the StratusLab distribution (or toolkit), a cloud middleware that integrates various software tools to provide customized infrastructure deployment capabilities. This middleware needs to be tested and validated in near-production environments to ensure that certain functional and Quality of Service requirements have been met. For this purpose one of the main tasks of WP5 is to setup and operate a reference cloud site using the software integrated by WP4. The operation of this site has been set as one of the main milestones for WP5 and was achieved in the first months of the project. The infrastructure is currently running release v0.3 – a beta release for the upcoming 1.0 production release.

2 Reference Cloud Service

2.1 Service Deployment Evolution

The reference cloud service follows the evolution of the software distribution provided by WP4. Work on the service started during the early stages of the project, before there was any cloud middleware available, with the installation of a plain OpenNebula site (version 2.0). This site was used for evaluation purposes as well as for demonstrations at the EGI Technical Forum 2010 in Amsterdam. The first publicly accessible site was provided with StratusLab release v0.1. A public announcement was distributed in various dissemination channels and the first requests for user accounts started arriving during the first hours of operations. During the following months the cloud site followed the evolution of StratusLab distribution upgrading to versions 0.2 and to 0.3, which is the latest version of the software and a beta release for the upcoming 1.0 production release. At the end of the first year of the project, the site will migrate to StratusLab v1.0.

2.2 Physical Configuration

The reference cloud service is currently hosted in one of GRNET's datacenters. Table 2.1 summarizes the physical resources allocated for the site. All the nodes are co-located in the same physical rack and are interconnected using 10 GigE (Figure 2.1). Storage is provided by a datacenter-wide storage server.

Table 2.1: Summary of the physical resources of the cloud site.

Number of nodes	17 (1 Frontend + 16 hosts)
Number of cores for VM instantiation	256
Total memory per physical node	48 GB
Total memory for the site	768 GB
Total storage for hosting VMs	3.6 TB shared over NFS
Number of public IPs available	120
Range of public IPs available for VMs	62.217.122.130–250
OS for hosts	CentOS 5.5

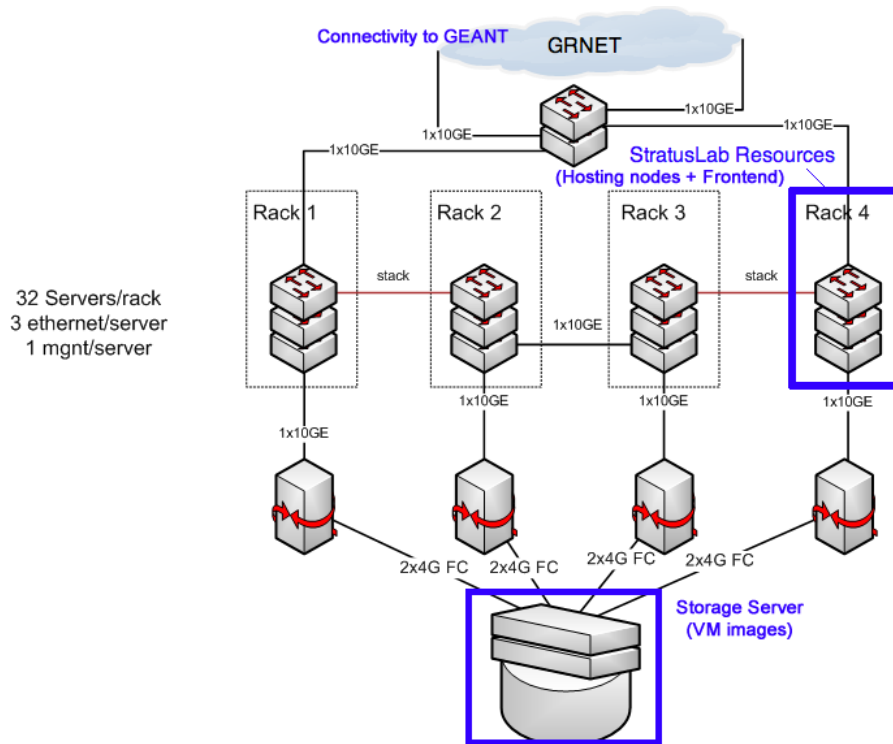


Figure 2.1: StratusLab's physical infrastructure in GRNET's datacenter.

Connectivity between the nodes and the storage server is achieved through dual 4 Gb/s Fibre Channel links. The above configuration enables the instantiation of a total of 120 VMs with public IP connectivity. Private or local IP interfaces are not yet available but they are planned for future upgrades of the service.

2.3 Service Provisioning

The reference cloud service is publicly available to everyone wanting to try out the StratusLab cloud middleware distribution. Users may request access to the service by sending an email to the support email list¹ providing their contact details and some basic information about their reasons for using the cloud. Until version 0.2 of the distribution the only option for accessing the service was with a username/password pair. As of version 0.3 users can access the cloud site via certificate based authentication. In the latter case, users must provide the DNs of their digital certificates, otherwise StratusLab will generate a username/password pair for them.

In order to access the service the users need to download and install the latest version of the StratusLab python command line tools available from the project's

¹support@stratuslab.eu



Enhancing Grid Infrastructures with
Virtualization and Cloud Technologies

Web Monitor

Nodes Instances

Enable auto refresh

List of nodes

Id	IP	Total CPU	Free CPU	Total mem	Free mem	Running VMs	State
0	62.217.120.155	1600	1598	49449772	46846440	3	Monitored
1	62.217.120.147	1600	1600	49449772	47011480	3	Monitored
2	62.217.120.148	1600	1596	49449772	47354468	2	Monitored
3	62.217.120.149	1600	1600	49449772	47250716	3	Monitored
4	62.217.120.150	1600	1598	49449772	47109348	1	Monitored
6	62.217.120.152	1600	1600	49449772	47011360	3	Monitored
7	62.217.120.153	1600	1600	49449772	48440180	2	Monitored
8	62.217.120.154	1600	1600	49449772	48233480	2	Monitored
9	62.217.120.156	1600	1600	49449772	48296108	2	Monitored
10	62.217.120.138	1600	1600	49449764	46574188	3	Monitored
11	62.217.120.139	1600	1598	49449764	46420368	3	Monitored
12	62.217.120.140	1600	1598	49449764	45764572	4	Monitored
13	62.217.120.141	1600	1596	49449772	44729112	4	Monitored
14	62.217.120.142	1600	1598	49449764	45473688	2	Monitored
15	62.217.120.143	1600	1600	49449764	46030140	3	Monitored
16	62.217.120.151	1600	1595	49449772	45717496	7	Monitored



StratusLab is co-funded by the European Community's
Seventh Framework Programme (Capacities)
Grant Agreement INF0-RI-261552



Figure 2.2: Web monitoring of cloud site (hosting nodes)

package repository² either in RPM (Red-Hat Package Manager), zip, or tarball formats. Using these tools and the service's endpoint³, the user has the ability to instantiate new VMs, inspect the status of existing running VMs and kill previously instantiated VMs. Information and a video tutorial for using the cloud service is available from the project web site⁴.

2.4 Site Monitoring

Two tools are used for monitoring the site. Ganglia⁵ has been setup for monitoring all physical hosts provided by GRNET for the purposes of the project. A subset of these hosts have been dedicated for hosting the cloud service and thus monitoring is available from Ganglia.

A web monitoring tool has been developed within the project by WP4 and is installed with every StratusLab cloud. The tool provides access to basic node and VM information. Figure 2.2 provides a snapshot of this tool depicting the total number of hosts that comprise the cloud site and the number of VMs running on them.

2.5 Sample Applications

During these months of the site's operations it has been used for various tests and for hosting a number of small and large applications. Most notably the site has

²<http://yum.stratuslab.eu>

³<https://cloud-grnet.stratuslab.eu:2634/RPC2>

⁴<http://www.stratuslab.eu>

⁵<http://node017.one.ypepth.grnet.gr/ganglia/>

been used for deploying a production grid site running gLite middleware [2] and for running Bioinformatics applications. The deployment of the grid site is reported in the next section. The details regarding the bioinformatics application have been reported in MS3 [1]. In brief, two virtual appliances, “Biological databases repository” and “Bioinformatics compute node”, have been prepared to provide bioinformaticians with a repository machine maintaining up-to-date international reference databases, then made available through shared filesystem in destination to bioinformatics cloud nodes with pre-installed bioinformatics software.

3 Hosting of Grid Services

During the first months we experimented extensively with the installation and operation of grid sites on top of cloud services using pre-configured VMs. At the end of this process we decided to formalize the installation of a virtualized grid site by certifying it within the EGI production infrastructure. In parallel we prepared a number of VM appliances for the basic node types of a gLite-based grid site: the Computing Element, the Storage Element, the Worker Node, the User Interface, and the APEL service. All of these images are available from the appliance repository¹. The appliances follow the evolution of gLite middleware—with every new release new image snapshots are created and uploaded to the repository.

The certified grid site is named HG-07-StratusLab. The site was certified within the GRNET NGI (the Greek National Grid Initiative) and has joined the national grid infrastructure (HellasGrid). The site offers a Computing Element and 8 dual-core Worker Nodes, thus providing a total capacity of 16 cores for job execution. The site also supports MPICH-2 and OpenMPI parallel applications. Communication among the nodes for MPI execution is supported through SSH host-based authentication. Each WN is configured with 4GB of main memory. Also, the site includes a Storage Element that offers a total storage space of 3TB. This storage is configured directly as an NFS mount-point from the local storage server and is not yet virtualized, as the cloud storage services of StratusLab are not yet available.

The site currently supports the StratusLab VO (vo.stratuslab.eu) as well as the required EGI-wide and local ops (operations) VOs. Obviously the job processing capacity of the site is currently rather limited. In this first phase the site primarily serves as a testbed for grid-cloud interoperability tests and for evaluating the implications of operating grid sites on public cloud services. Already a number of issues have been identified which will be reported in the coming quarter. In the future depending on the workload and potential requests to support additional VOs, it should be rather trivial to expand the workload execution capacity of the site (i.e. number of available cores and/or Worker Nodes).

The GStat page with all the details of the site as they are reported from the Site-BDII are available².

¹<http://appliances.stratuslab.eu>

²<http://gstat-prod.cern.ch/gstat/site/HG-07-StratusLab/>

References

- [1] C. Blanchet. Creation of Virtual Appliances for Bioinformatics Community, 2011. StratusLab project milestone MS3.
- [2] gLite. Lightweight Middleware for Grid Computing. <http://glite.web.cern.ch>.